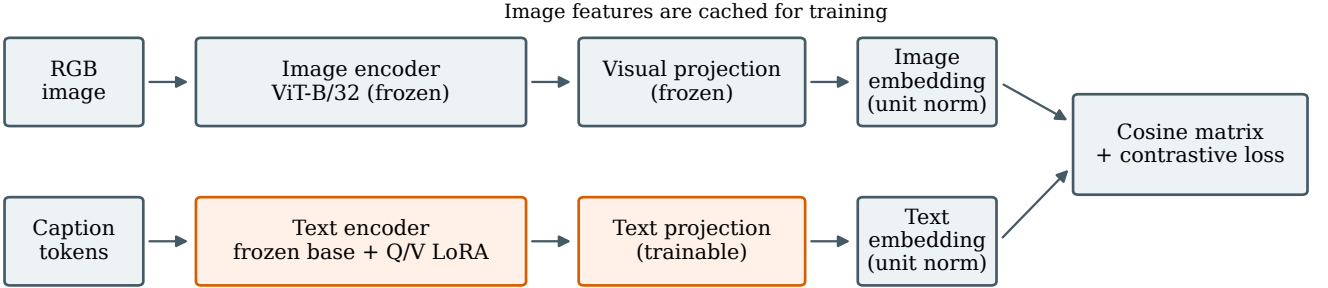
Image features are cached for training
RGB image
Image encoder ViT-B/32 (frozen)
Visual projection (frozen)
Image embedding (unit norm)
Cosine matrix + contrastive loss
Caption tokens
Text encoder frozen base + Q/V LoRA
Text projection (trainable)
Text embedding (unit norm)
Trainable in C: 24 text Q/V LoRA modules and text projection